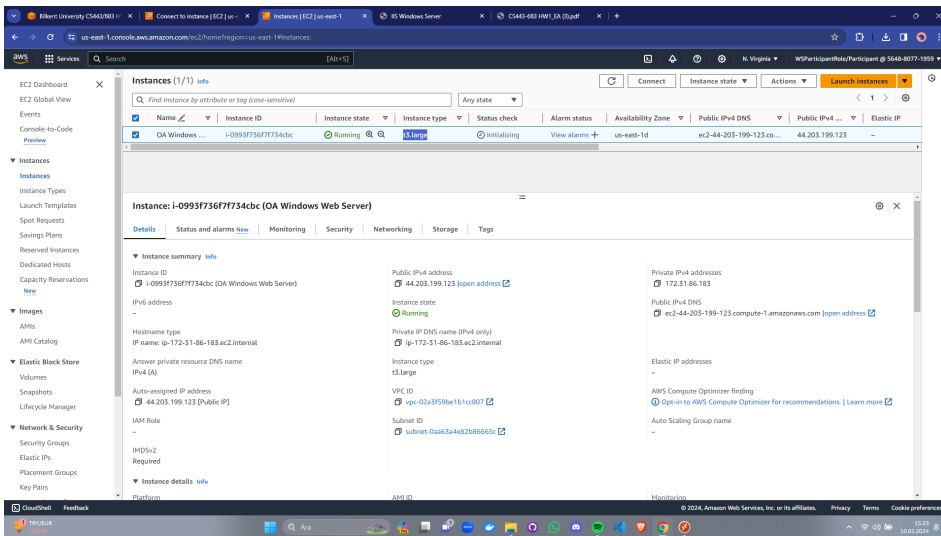
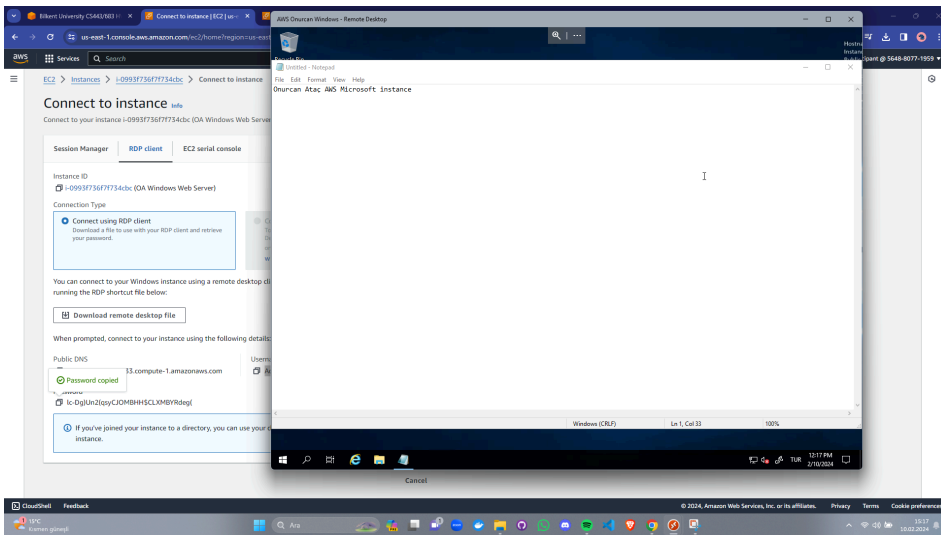
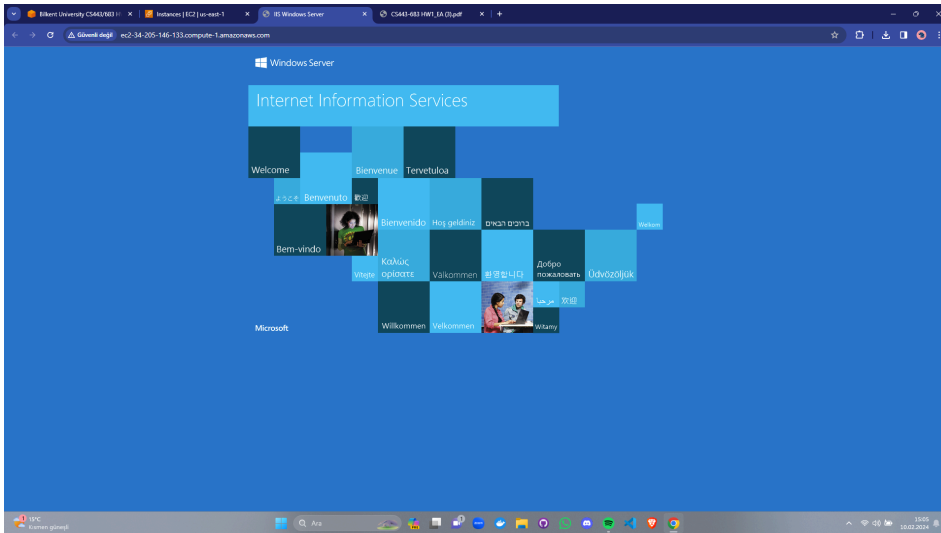
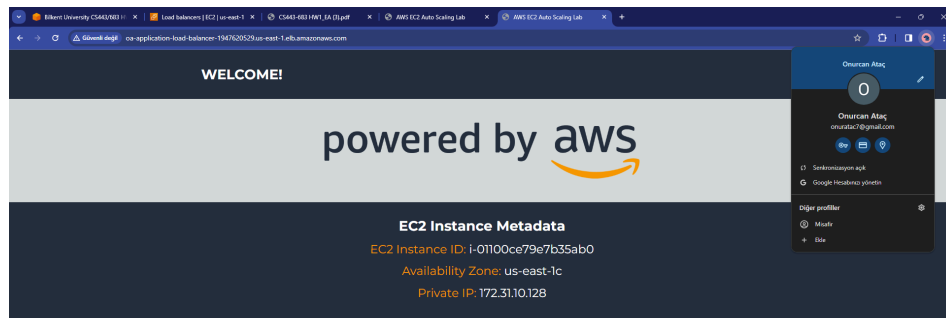


In the second part of the assignment, I created an EC2 instance that runs Windows. The process was similar to the first part. I used the AWS Console. I created a new key pair, and launched a new instance using Windows quick start AMI. I then connected to the instance that I created using an RDP client, in my case I used Microsoft Remote Desktop application. I then stopped the instance and changed the instance type from t2.medium to t3.large easily. I also created an elastic IP for my server. I cleaned up the resources to finish this part.



3. Auto Scaling on AWS

In the third part of the assignment, I learned how to create an auto scaling group that scales and adjusts automatically when the demand increases. I started by creating a web host by using the “.yaml” file the lab provides and AWS CloudFormation service. I then created a security group. I then created a launch template that specifies the launch parameters to enable an easy launch for the instances while automatically scaling. I then set an auto scaling group up to determine the default, minimum and maximum number of instances. I then created a load balancer security group, in order to specify inbound and outbound rules, which enables the load balancer to send its outgoing traffic only to the specified auto scaling group. I then increased the CPU load and triggered the creation of new instances automatically. I cleaned up the resources and completed the assignment.



Current CPU Load: 0%

